

LLM-use in biomedical publications - an estimation

Lena Holzwarth

March 2026

1 Introduction

Since the release of OpenAI’s ChatGPT¹ in November 2022, chatbots based on Large language models (LLMs) have become ubiquitous in the public consciousness. Due to their continued improvement, going from simple chat interfaces to fully integrated agentic systems like Anthropic’s Claude Code², they are increasingly easy to embed in any workplace dealing with writing, programming, or project planning. One such workplace is academia, and it is not surprising that researchers across disciplines have begun to adopt LLMs for editing or writing papers, peer reviews, or grant applications[30]. Proponents of the technology praise it as a tool for accessibility and inclusivity[5, 6, 52], arguing that it enables non-native English speakers to overcome language barriers by translating or performing grammar checks[29]. However, the surge of ChatGPT’s popularity is accompanied by concerns of misuse[37, 61]. Because LLMs are prone to hallucination, if their output is not thoroughly checked before using it in a paper, false or misleading statements might be published, especially if the review is also done by an LLM[2, 39, 57, 60]. Chatbots also make scientific fraud much easier, with the possibility to fabricate entire papers at the cost of a single prompt[26]. Reviewers relying on LLMs to generate their review are vulnerable to hidden prompts in the paper’s text, telling the LLM to give a positive review[38]. In light of these concerns, many journals have issued guidelines about the disclosure of LLM-use[8]. However, the perception of LLM-use leading to inaccuracies in the text, together with journals restricting or explicitly banning certain use-cases[54] may incentivize academics against disclosing. Consequently, when trying to find out how prevalent LLM-use is in academia, and especially in the paper-writing process, relying on self-disclosure statements will lead to underestimation[28, 30].

Regardless of the normative value assigned to LLMs in their function as paper-writing assistants, their prolonged use in academic text production will have an impact on academic writing as a whole. LLMs have a distinct writing style. Sentences generated by them have less variability in length[10], but higher lexical diversity and positive sentiment[42]. Famously, they also show a preference for certain words, an observation which has led to much ridicule around the most popular example – let’s delve deeper[41]. The reason for this behavior is not entirely clear. There has been no evidence so far to suggest that the training data, model architecture, or training algorithm have an influence in forming these idiosyncratic word choices[24]. However, the authors were able to induce an LLM with preferred words during reinforcement learning with human feedback (RLHF)[4], the continued training scheme behind the success of modern Chatbots. During RLHF, humans provide learning goals by rating the interaction with the model. This work is largely performed by underpaid workers in the global south[50], for example in Nigeria, where, incidentally, “delve” is used much more frequently than in other English-speaking countries[23].

These selective word preferences can be leveraged to estimate the proportion

¹<https://chatgpt.com/>

²<https://claude.com/product/claude-code>

of LLM-processed text in a corpus. [27] track changes in word frequencies in biomedical abstracts from the PubMed database, by comparing a word’s observed frequency in a given year with a projection based on its frequencies two and three years back. The metric obtained by subtracting the predicted from the observed frequency is called excess frequency: A change in frequency that diverges from previous trends. They identify many words that appear with excess frequency in abstracts from 2024. This is notable, because for previous years, the observed frequency gaps are much smaller. It’s possible to apply a threshold to excess frequency, such that most words from pre-2024 don’t exceed it, but many from 2024 do. In addition, the majority of pre-2024 excess words are clearly related to a specific research topic, such as “ebola” in 2015, “convolutional” in 2017, or “coronavirus” in 2020. In contrast, many excess words in 2024 are content-independent style words – “underscore”, “potential”, “significant” or “pivotal”, to name just a few. The stark difference between excess words from before and after 2024 point to a drastic shift in writing patterns that can only be explained by the adoption of LLMs into the writing process. All identified excess style words for 2024 can therefore be interpreted as marker words for LLM use. It’s important to note that the marker words apply to analyses on the corpus level and not for individual documents – one “delve” does not an LLM-generated text make. While these words appear more frequently in LLM outputs, some of them still have very low frequency in both human and machine text. There are texts by human authors that include them, and LLM texts that don’t.

Using the combined excess frequencies of a selection of excess style words, [27] derive a lower bound for LLM use estimates. Their reasoning is quite intuitive: all instances of using one of the relevant words that go beyond the predicted frequency must have been caused by LLM use. Excess frequency as a metric has the advantage of not needing to distinguish between using an LLM write a whole paragraph, or suggesting changes on previously written text. Both use cases will lead to some LLM-suggested words being present in the final product.

Here, we expand on their work on excess word frequency as a lower bound estimate of LLM use in biomedical abstracts by making the following contributions:

- **Widening the focus** from an analyzing only the abstract to the entire paper. By using the open-access PubMed Central (PMC) dataset[44] instead of PubMed, the proportion of LLM-assisted text in the main text of the paper can be compared with usage in the abstract, and different sections within the body can also be compared against each other.
- Deriving a **stronger lower bound** metric going beyond excess frequency. While excess frequency is an intuitive estimate for LLM use, it severely underestimates the true proportion of LLM-assisted text, as will be shown.
- Computing **uncertainty measures** and performing tests on simulated data to increase confidence in the estimates. A disadvantage of the previous approach is the lack of a ground-truth set to test performance with.

By simulating the data generation process in a Binomial model of word occurrence, we can evaluate our metric for several ground-truth levels of LLM use.

- **Comparing** the LLM use estimates with those of other approaches at estimation, as well as survey data asking researchers about their LLM use. Unlike previous studies, we find that LLM-like language has almost completely permeated academic writing.

This thesis is structured in the following way: In Section 2, other approaches to estimate LLM-use in academia are discussed, such as comparable word frequency measures, word distribution modeling and LLM-detectors. Section 3 gives an overview of the data processing pipeline, the metrics used to estimate LLM-use, as well as the validation procedure employed to increase confidence in the estimates attained. This is followed by Section 4, where our estimates are reported and compared to results from other work. Finally, the results, along with their impact and limitations are discussed in Section 5.

All code is available on GitHub³. The PMC dataset can be downloaded here⁴.

³https://github.com/LenaHolzwarth/llms_in_academia

⁴https://ftp.ncbi.nlm.nih.gov/pub/pmc/deprecated/oa_bulk/

2 Related Work

Since the wide-spread proliferation of LLMs, many techniques to identify traces of LLM-use have been suggested. As this work is focused on LLM use in academia, the approaches detailed here have been selected accordingly with no claim of exhaustiveness.

2.1 Word frequency measures

Most closely related to our approach, approaches working with word frequency measures make use of the distinct word choices separating LLM and human text production. While not immediately visible in individual samples, these patterns emerge when looking at word distributions on the corpus level. Reports of certain words appearing with disproportionate frequency after the introduction of ChatGPT in November 2022 have been made for publications from a wide array of scientific disciplines, ranging from linguistics[7] over astronomy[3], to medicine[43] or dentistry[56].

These observations of increased word frequency can then be used to derive estimates for the prevalence of LLM use in the respective databases, using the words with the most pronounced frequency changes as markers of LLM use. For example, [19] computed the relative yearly change in word frequency in papers from the Dimensions database, both for single words and word groups of up to twelve terms, which were selected from a list of LLM marker words identified by [34]. Their LLM-use estimate is based the frequency change of these excess words / word groups. The word frequencies of 2022 are multiplied with the highest frequency increase between two consecutive pre-ChatGPT years to obtain a conservatively estimated counterfactual baseline of human word frequency. The number of papers with at least one of the marker words in 2023 exceeding the amount explained for by the baseline divided by the total number of papers in 2023 servers as their LLM estimate. A later study from the same author [20] repeated the same procedure with an updated word list, following reports that some marker words showed decreased usage in 2024 [16], making them less useful as LLM-use markers.

[29] employed a similar approach of monitoring frequency changes of marker words in abstracts from Scopus, Web of Science and PubMed, as well as full papers from PMC, Dimensions and OpenAlex. They selected some of the marker words reported in previous research, followed by collecting additional marker words by searching for words with high co-occurrence with the pre-selected list. While they observe strong increases in marker word frequency between the years 2022 and 2024, they don't use this to estimate LLM-prevalence in their data. However, based on their results, [20] derives an estimate by comparing the frequency of the word "underscore" in observed 2022 directly to its frequency in 2024, not taking into account that human usage of the word is also subject to yearly changes.

2.2 Word distribution modeling

Similarly to the approaches detailed in the previous section, word distribution modeling makes use of distinctive word frequency patterns in LLM-generated text. However, instead of relying only on observed frequencies in the target corpus, enabling the estimate of a ground-truth of word frequencies in human text when limiting data to pre-2023, it goes one step further by simulating LLM text to also obtain a baseline estimate for LLM word frequencies.

Using a technique they call “distributional GPT quantification”, [34] analyze peer reviews of AI conference papers. Their work is based on observation that word distributions are different between LLM and human-written text, and the documents are generated from a mixture distribution of human and LLM text, the fraction of LLM text in the sample can be estimated via maximum likelihood estimation on the corpus, if the probability distributions for human and LLM-generated text are known. The probability distribution for human text can be estimated by counting word occurrence in texts from pre-LLM years. To obtain the probability distribution of words in LLM-generated text, the authors prompt ChatGPT to generate alternative reviews of the conference papers from before 2023. The authors modeled the distribution of all adjectives instead of relying only on marker words associated with LLM use. As a validation of their method, a corpus with known LLM-proportion was constructed by sampling from the original and generated reviews. By computing the MLE for this corpus, the researchers were able to retrieve the correct proportion, while using state-of-the-art AI text detection methods yielded much worse results.

The authors applied the same method to abstracts and introductions of scientific papers sourced from arXiv, bioRxiv and Nature portfolio journals[33, 35]. LLM text was sampled by tasking ChatGPT to create a summarization outline of a paragraph from an existing, pre-ChatGPT paper, which was then given to the LLM with a prompt to generate a full paragraph based on this outline.

A slightly different method was introduced by [17], who investigated LLM-use in arXiv abstracts. Like the previous approach, they based the human word frequency baseline on pre-2023 papers and LLM frequency baseline on outputs from ChatGPT asked to polish or re-write pre-2023 abstracts. Then, they compute a word’s frequency change rate by comparing its frequencies in the original and rewritten abstracts. The authors argue that estimates derived from a word list are more robust than those derived from single words, and further claim that the choice of the word list yielding the most accurate estimate depends on the (unknown) underlying value of LLM-use. To find the optimal word list for each of a set of ground-truth LLM-use values and validate their procedure, they create test data with a fixed proportion of original and LLM-edited abstracts. For each ground-truth proportion value, a word list is created by optimizing thresholds on word frequency and frequency change rate. Combining the observed frequency with the simulated frequency change rate, they can derive estimates for LLM use in the real data using all optimized word lists. The mean of the estimates is taken as the preliminary result. The estimates are then repeated

with the three word lists optimized for the LLM-use proportions closest to the preliminary result.

The validation procedure with ground-truth LLM-generated text is a marked advantage of these studies, compared to the frequency-based estimates, where no validation set with known distribution is available. However, this validation is strongly dependent upon the exact model and prompt that were used to create the LLM corpus, and the estimated distribution based on this corpus might differ drastically from the true LLM distribution created from a multitude of prompts and models. To illustrate this point, [17] perform validations on papers edited with a minimally different prompt than the one used to derive the frequency change rate, which leads to a systematic overestimation of true LLM use. While the true value is still within the error bounds, stronger alterations to the prompt or the use of a different model are likely to have a stronger impact on the estimates.

2.3 LLM-detectors

A more fine-grained approach to observing or modeling word distributions across entire corpi is to train classifiers on identifying LLM-use in individual texts.

Zero-shot LLM detection or model self detection uses properties of the model’s idiosyncratic text generation process to determine how likely it is for a text sample to have been generated by the LLM. An early example of this approach is **GLTR** [15], which determines for every word in the text which ranking it has among the LLMs suggestions given its position in the text. The lower the average ranking of the selected words, the more likely it is for the text to have been written by a human author. **Detect-GPT** [46] uses a separate LLM to generate perturbed versions of the sample and computes the log-likelihood of the original and perturbed versions under the target model, building on the observation that the log-likelihood of a text sample generated by the target LLM will be higher than that of the perturbed variants, while the same is not true for human-written text. **Binoculars** [22] uses a similar approach by normalizing the perplexity of a sample text under one model by the cross-perplexity of the re-generation of the sample text by a different model. The perplexity is the negative log-likelihood of a text and can be thought of as the model’s surprise, i.e. how far the text diverges from something the model would have produced. Binoculars has been used to check papers from the preprint platforms arXive, bioRxiv and medRxiv[9]. The disadvantage of these approaches is two-fold: First, in order to classify a piece of text, one needs to know with which LLM it was potentially generated, and second, one needs access to the scoring function of the LLM. While this is possible for some GPT variants, the scores of the most widely used model, ChatGPT, are not open to the public[46]. Also, these methods assume the whole text to be LLM generated, which would impede their accuracy on samples that were only partly edited with LLMs.

Trained LLM detectors, on the other hand, rely on a training corpus of ground-truth LLM and human text. As an alternative to zero-shot detection, they circumvent the need for access to model parameters and can, in

theory, work model-agnostically, if the training set is constructed with samples from different models. Their performance is, however, critically dependent on the prompts and models used to generate the training corpus[40]. A common approach is to fine-tune a pretrained language model to distinguish between LLM-generated and human-written text. Another possibility is to train conventional classifiers such as XGBoost or support vector machines on embedded representations of training text[5]. **Feature-based detectors** use linguistic features, such as the variation in sentence length as indicators of LLM use[10]. The classifier is then trained not on the texts themselves but on the features extracted from them.

There is also a plethora of proprietary LLM detectors on the market, the majority of which do not disclose their model architecture or training procedure. Because of their relative convenience, many researchers investigating LLM use in academia opted for these models to generate estimates[1, 40, 51]. Among the models used are originality.AI, GPTzero, zeroGPT, GPTKit, Smodin, Sapling, Content at scale AI detector and Writer AI content detector.

Many concerns have been raised about current LLM-detectors' lack of accuracy[11, 31, 58], going so far as to dispute even the possibility of reliable LLM-detectors[53]. Early methods have been found to discriminate against non-native English speakers, disproportionately flagging their work as AI-generated, and to be unable to adjust to slight prompt changes[32].

2.4 Surveying scientists

Perhaps the most straightforward approach to evaluating the prevalence of LLM-use in academia is to ask the researchers themselves. Several surveys on scientists' use of AI have been conducted by publishing houses[12, 30, 49, 48, 59] and independent researchers[13, 36, 45, 47]. These studies can be hard to compare because of different sampling pools and phrasing of questions. An overview of questions asked in each of the surveys can be found in the Appendix. While some researchers also acknowledge the use of LLMs in their papers, the number of papers crediting ChatGPT or similar models is much smaller than even conservative estimates of actual LLM use[28], emphasizing the need for anonymous surveys to produce a more accurate picture of researchers' attitudes towards the technology. For example, in a survey done by Nature, 28% of respondents reported to have used LLMs to edit their papers, but only 10% disclosed their AI use[30].

3 Methods

3.1 Dataset and Preprocessing




Figure 1

Pubmed Central

Pubmed Central (PMC) is a collection of open-access biomedical literature curated by the U.S. National Institutes of Health’s National Library of Medicine [44]. It provides machine-readable versions of the articles in XML format, which can be accessed individually or downloaded in bulk. There are three open-access subsets available for bulk download with different licensing: commercial, non-commercial and other. The latter set includes non-machine readable licenses and was therefore not used. PMC publishes a baseline twice a year containing all papers in the collection up to shortly before the baseline’s release date. Incremental updates are published daily in separate files. When a new baseline is released, the previous baseline and all incremental files are replaced. Here, we used the baseline published on January 23, 2026 containing papers dated December 31, 2025 and before.⁵ The daily updates published after the baseline were not considered.

XML parsing

From the XML files in the baseline, each corresponding to one paper, the following information was extracted:

- Article type: specifies whether the article is a research paper, a review, a comment, etc.

⁵Since work on this thesis started earlier, initial analyses used the previous baseline released June 26, 2025. These were repeated with the January 2026 baseline, so that everything reported here is based on the newer baseline

- Date
- Language: only specified for a minority of papers, but is inferred later during preprocessing
- Country
- PMC-ID: unique identifier for every paper in PMC
- PMID: unique identifier for PubMed, not available for every paper in PMC
- Journal name
- Title
- Abstract
- Sections
- Section titles

The sections and section titles are parsed from the main body of the XML file. In general, individual sections are tagged as `sec` and assigned a title. This is important for being able to compare different sections with each other, for which we need as many papers as possible with similar sections. However, since there is no standardized format for the files, some sections are untitled or there are no section tags at all. In particular, many files have section names for all sections except the Introduction, or don't embed the Introduction in a separate section tag. Therefore, if a file has more sections than section titles, and the first section is unnamed, it is assumed to be the Introduction. Any other untitled sections are discarded. Similarly, if there is text at the beginning of the body outside of a `sec` tag, it is assumed to be the introduction. This heuristic is not ideal, because the unnamed text contains abbreviations or keywords instead of the introduction in some cases, but since the Introduction is among the sections this study wants to focus on, it strongly increases the number of usable papers.

Preprocessing

While some XML files specify the paper's language, this is not the case for the majority. Since the focus here is on English text, we determine each paper's language using `polyglot`'s detector module. Every non-English paper is discarded. The dates of each paper are standardized to DD-MM-YYYY format. Some entries only specify the month (e.g. "Feb 2025") or season of the year (e.g. "Fall 2025"), in which case the first of the month or start of the (meteorological) season are used as dates. Papers without date are removed. Since we want to compare LLM-use across different sections, it is important to have as many papers with comparable section structure as possible. To achieve this, in a first step the section names were manually standardized by collecting the most frequent section names and mapping them to their standardized version. For example, *experimental results*, *findings* and *main results* would be mapped to

results. If a paper has multiple sections with the same standardized title, they are concatenated. For example, *data* and *experimental methods* would first both be renamed as *methods* and then combined into a single *methods* section. Originally, we planned to investigate different article types and article structures, but in the end we focused only on research articles that have exactly the sections Introduction, Methods, Results, and Discussion, since papers with this article structure were the majority in the dataset. Conversely, all papers with article types other than *research-article* or with more, fewer or different sections were discarded. Finally, all sections, including the abstract, were required to have a length of at least 250 characters. These preprocessing steps reduce the dataset from 7,125,722 to 1,602,404 papers, 1,194,287 of which are dated 2017 or later. The use of PMC, granting us access to the whole paper, enables an extension on [27] in two regards: First, we can compare LLM-usage in the abstract to LLM-usage in the entire body of the paper, and second, we can compare different sections within the body with each other. For this purpose, all of the analyses detailed in the following subsections were performed on the abstract, the introduction, the methods, the results and discussion of the paper, as well as the full paper, which refers to the four sections from the body without the abstract.

3.2 Excess Frequency and LLM-use lower bound estimates




Figure 2: Shared caption describing both images

Excess frequency, referring to the difference between the observed and predicted frequencies for some word w , can be thought of as an intuitive lower bound for LLM use. To compute it, we first determine the observed frequency of w . We use sklearn’s binary CountVectorizer to count the number of papers in each month or year in which w occurs. Because it is a binary count, it doesn’t matter how often w appears in each paper. The vectorizer only considers words with 4 or more letters, all of which must be part of the English alphabet. When counting word occurrence, lemmatization is an important consideration: one has to determine whether to treat different inflections of w as the same or as individual words. With lemmatization, it makes no difference if a paper contains *delves*, *delved* or *delving*, they are all treated as an instance of *delve*. It

is not immediately clear if lemmatization is helpful in this use-case. We know from previous studies that LLMs prefer certain words. This could imply that LLMs use all inflections of a word at a higher frequency, or that it has a strong preference for that word in a specific grammatical context (e.g. third-person singular), while using every other inflection at the same frequency in which they occur in human-written text. As will be detailed in Section 3.3, for the overall LLM-usage estimates, this does not matter. For the plots showing results for individual words, such as Figure, all word variants are included in the counts. The frequencies are computed by dividing the monthly or yearly counts by the total number of papers in each respective time frame. On a conceptual level, the observed frequencies q are an estimate for the probability of a positive outcome in the Bernoulli trial of whether a word is present in some paper. For dates prior to the release of ChatGPT in November 2022, marking the beginning of easily accessible LLMs, q corresponds to the frequency of a word in human-written text, p_{human} . After the widespread availability of Chatbots, the observed word frequency can be thought of as the product of a mixture distribution between the probability of a word to appear in human-authored and LLM-assisted text, respectively:

$$q = \begin{cases} p_{human} & \text{for papers before 11-2022} \\ (1 - \beta) \cdot p_{human} + \beta \cdot p_{LLM} & \text{for papers after 11-2022} \end{cases} \quad (1)$$

The mixture component β in the lower case indicates the proportion of papers that were written with the help of LLMs. We can transform this equation to obtain an expression for β :

$$\begin{aligned} q &= (1 - \beta) \cdot p_{human} + \beta \cdot p_{LLM} \\ \beta \cdot (p_{LLM} - p_{human}) &= q - p_{human} \\ \beta &= \frac{q - p_{human}}{p_{LLM} - p_{human}} \end{aligned} \quad (2)$$

While p_{human} can be estimated by the observed q in the time span before November 2022, there is no direct way to do so after the release of ChatGPT. However, if we assume that human usage of words follows linear trends and is independent of LLM-use, we can use the values before 2022 to model $\hat{p}_{human}$ using linear regression. Since we are interested in both the monthly and yearly trends, we fit the regression either on the five yearly frequency values from 2018 to (including) 2022, or on the 60 monthly frequencies from November 2017 to October 2022. For the yearly data we include 2022 in the train set because, considering the time it takes to write a paper, and the time it took the new technology to gain traction, it is unlikely that many papers published in December 2022 were written with the help of LLMs. Replacing the unknown p_{human} with the estimate $\hat{p}_{human}$ in equation 2, one issue remains: there is no direct way to estimate p_{LLM} from the observed frequencies. The others solve this problem by using ground-truth LLM-generated text to determine p_{LLM} .

Here, we choose a different approach by determining a lower bound for β that doesn't rely on p_{LLM} :

$$\begin{aligned}\beta &= \frac{q - \hat{p}_{human}}{p_{LLM} - \hat{p}_{human}} \\ \beta &\geq q - \hat{p}_{human} = \Delta\end{aligned}\tag{3}$$

This is the excess frequency measure introduced by [27] and denoted Δ . Intuitively, if a word appears in 70% of papers in 2025, but we predict humans to use this word in only 60% of instances, the 10% gap must be accounted for as having been written by LLMs. Note that this metric will lead to a strong under-estimation of the actual LLM usage, because a human-use frequency of 60% will only translate to 60% of observed use, if all texts in the sample are in fact written by humans ($\beta = 0$). In reality, some papers were written with the help of LLMs, such that the fraction of word-occurrence that can be explained by human usage is not 60% of all papers, but 60% of the papers with human authorship. A stronger lower bound, denoted $\hat{\beta}$ is also possible:

$$\begin{aligned}\beta &= \frac{q - \hat{p}_{human}}{p_{LLM} - \hat{p}_{human}} \\ \beta &\geq \frac{q - \hat{p}_{human}}{1 - \hat{p}_{human}} = \hat{\beta}\end{aligned}\tag{4}$$

In dividing by $1 - \hat{p}_{human}$, this lower bound readjusts the estimate by focusing only on the fraction of papers for which no human use of the word is to be expected. In the above example with $p_{human} = 60\%$ and excess frequency of $\Delta = 10\%$, $1 - \hat{p}_{human} = 40\%$ of papers wouldn't contain the word if there were no LLM-authored texts, but only $1 - q = 30\%$ actually don't, so the remaining $\frac{10\%}{40\%} = 25\%$ must be LLM-authored.




Figure 3

3.3 Grouped word frequencies




Figure 4: Fix title for delve estimates!

While single words are useful to illustrate the concept of excess frequency, they are suboptimal as a basis for an accurate estimate of overall LLM-use. The reason for this is obvious: while LLMs might use a word such as “potentially” at a higher frequency than humans would, “potentially” doesn’t fit in every context so there will potentially be a lot of LLM texts not containing it. Also, the observed frequencies of some words might be unreliable. The word “delve”, used as the titular example of excess vocabulary in [27], has experienced a sharp decline in popularity starting around mid-2024, when it was widely reported as an example of “LLM-speak”. Many authors have thus moved to using word lists as markers. The method of computing the frequencies doesn’t change, but instead of counting in how many papers a single word appears, we count in how many papers at least one of the words in the list appears. The most systematic way to create a list of words indicative of LLM-use is to use a data-driven approach, which doesn’t require a manually selected list.

Excess words for 2024 PubMed

[27] determine the excess words for 2024 by computing the frequency gap and the frequency ratio for each word in their dataset. The frequency gap is the same metric as Δ in Equation 3, although they only use the yearly frequencies of 2021 and 2022 to predict $\hat{p}_{human}$. The frequency ratio is computed as $\frac{q}{\hat{p}_{human}}$. An excess word is defined as any word with frequency gap or frequency

ratio above a certain threshold. The thresholds were chosen in such a way that the majority of words in texts from the years 2012 - 2023 were excluded (The gaps for previous years were also computed by comparing observed values to $\hat{p}_{human}$ computed based on two and three years prior). Using both metrics as inclusion criteria serves to increase the number of excess words. For instance, a word whose usage increases from 0.0001 to 0.001 will have a relatively small gap but a large ratio, while a word whose prior frequency of 0.3 rises to 0.4 has a large gap with a relatively small ratio. Both changes mark a notable difference in usage and should therefore be candidates for marker words. All words above the threshold were then manually annotated as *style words* or *content words*. *Content words* are all words related to the research topic of a certain paper. Their excess use can be attributed to new inventions or research trends. Examples include “RNN” in 2018, “Covid” in 2020, and “ChatGPT” in 2024. *Style words*, on the other hand, are not related to research topics and could be found in any discipline: “these”, “potentially”, “delve”, etc. Excess words of 2024 were defined as all style words with frequency gap or ratio gap above the respective thresholds. This resulted in a list of 379 words, which can be found in the Appendix B.

Optimal word lists

The excess word list can be used to create the optimal list of marker words to detect LLM-use. It might be counter-intuitive to not just use the whole list: if any of the words can act as a marker for LLM-use, then using as many words as possible should yield the most accurate result. However, one has to keep in mind that all excess words also appear in human-written text. If there are too many words in the list, their cumulative frequency will be close to 1 even for p_{human} . For [27], who use the frequency gap Δ as their LLM-use estimate, this is suboptimal because the closer $\hat{p}_{human}$ is to 1, the smaller the potential gap. When using $\hat{\beta}$ as the metric, the best estimate would be achieved by getting p_{LLM} as close to 1 as possible, which would in turn move $\hat{\beta}$ closer to the true value β . While p_{LLM} isn’t observable, it can be manipulated by increasing the number of words in the excess list. However, this also includes the caveat that $\hat{p}_{human}$ must remain smaller than 1, otherwise, Equation 4 yields division by zero. This means that for both estimates, Δ and $\hat{\beta}$, a subset of the 2024 excess words needs to be determined. Because brute-forcing the decision on an optimal list by trying all possible combinations is not feasible, [27] sorted the excess words by their individual frequencies in 2024 (lowest to highest), and then used different cutoff values to create subsets of all words whose individual frequencies were below the cutoff. Each resulting list was used to compute the LLM-use estimate Δ and the highest Δ was used as the final estimate.

Adaptation to PMC data Initially, we planned to repeat the procedure of finding excess words for the PMC dataset. This turned out to be non-feasible, because the length of the full papers, as opposed to only abstracts, lead to a huge number of excess words, which would all have to be manually annotated to filter them for style words. Instead, we used the previous study’s 2024 excess word list.

This is permissible because, first, PMC is a subset of the PubMed data used in [27], so the domain is similar, and second, the point of using style words is to have marker words that are domain-independent. (Although it is likely that when searching for excess words in a different domain, such as news articles, the resulting list would be somewhat different). The 2024 excess word list is not lemmatized, so that “delving” and “delves” are treated as different words. Since we are not interested in the excess words on a semantic level and use them only as indicator of LLM use, it does not matter if two words on the list are variations of the same verb, or if a third variation of the verb is not included. To get the optimal list for each section, we sort the 2024 excess words by their individual frequency in 2024. This leads to a somewhat different order for each section. Like in the previous work, we then use different cutoff values to create subsets of excess words, which serve as a basis for computing the LLM-use estimate $\hat{\beta}$. We started with x cutoffs spaced evenly on a log-scale between 0 and 1 and later added intermediate values in the high-interest region around the maximum, resulting in y cutoffs. Initially, we used the same approach as above to select the maximum $\hat{\beta}$ -value as the estimate. However, in some cases, the maximum was at one of the highest cutoff value, leading to nearly the full excess word list being used for the estimate. This results in values for q and $\hat{p}_{human}$ that are extremely close to 1, sometimes with $\hat{p}_{human}$ being larger than q , which contradicts the assumption $p_{human} < p_{LLM}$. As a consequence, both numerator and denominator of the fraction $\frac{q - \hat{p}_{human}}{1 - \hat{p}_{human}}$ are very close to 0, which can lead to unstable behavior. To prevent this, any cutoff that leads to estimates of $\hat{p}_{human} \geq 0.999$ was excluded as candidate for creating optimal excess word lists.

The optimization was performed on yearly frequency data separately. This means that the optimal cutoff for 2024 could be different to that of 2025. For the main analysis, the optimal cutoff values for 2025 were used for all time points to simplify procedure and have more easily comparable results. Notably, the main analysis was performed on monthly data with the word lists optimized for the year 2025, instead of optimizing each month individually, for simplicity. A possible approach with optimizing on monthly intervals of the data would be to take the majority optimum as the list for all estimates. The end result would be the same: Some time intervals would be estimated with the optimal word list, while others wouldn’t.

To better judge the reliability of the estimates, we computed the standard error and performed sanity checks using simulated frequency data.

3.4 Method validation

Standard error for LLM-use estimates Since the LLM-use estimate $\hat{\beta} = \frac{q - \hat{p}_{human}}{1 - \hat{p}_{human}}$ depends on $\hat{p}_{human}$ and q , we first need to determine their respective errors, before propagating them to obtain the error for $\hat{\beta}$. We conceptualize q as the maximum likelihood estimate (MLE) for the binomial probability of a word from the excess word list to appear in a paper. As such, the MLE’s variance is $SE_q^2 = \frac{q(1-q)}{n}$, with n as the number of papers in the current month or year.

$\hat{p}_{human}$ is the regression estimate of a binomial probability. Since $\hat{p}_{human}$ is not an empirical estimate, the binomial MLE variance doesn't apply and we focus on the regression variance. Technically, the data points that the regression was trained on are empirical estimates (since training was done on q values in the time span before the release of ChatGPT). To factor in this binomial maximum likelihood uncertainty, it should be included in the regression formula. However, because the MLE variance is very small due to the large number of papers each month, this was omitted for simplicity. The regression variance for design matrix X (including only points on which the regression is fit, corresponding to the time before the release of ChatGPT), predictor matrix x (including all points that are to be predicted, including pre- and post-ChatGPT frequencies) and response vector y is computed with the following formula:

$$SE_{\hat{p}_{human}}^2 = \hat{\sigma}^2 x(X^T X)^{-1} x^T + \hat{\sigma}^2 \quad (5)$$

with $\hat{b} = (X^T X)^{-1} X^T y$ as the estimated regression coefficients⁶ and $\hat{\sigma}^2 = \frac{\|y - X\hat{b}\|^2}{n-2}$ the unbiased estimate of the noise variance, where n refers to the number of papers, of which the number of predictors, 2, is subtracted. $\hat{\sigma}^2 x(X^T X)^{-1} x^T$ is the covariance matrix of the regression coefficients, specifying the uncertainty of the mean prediction. To this, an additional $\hat{\sigma}$ -term is added to estimate how far the predicted values would have deviated from this mean.

The variance of $\hat{\beta}$, as a nonlinear function of q and $\hat{p}_{human}$, is not trivial to determine. We can apply the Delta method, which uses Taylor series approximation, to do so. We assume q and $\hat{p}_{human}$ to be the means of random variables Q and P , respectively. First, we approximate $\hat{\beta}((Q, P))$ around the mean $\mu = (q, \hat{p}_{human})$ with a first-order Taylor expansion:

$$\hat{\beta}((Q, P)) \approx \hat{\beta}(\mu) + \nabla \hat{\beta}(\mu)^T ((Q, P) - \mu) \quad (6)$$

with gradient ∇ as

$$\nabla \hat{\beta}(\mu) = \left[\frac{\partial \hat{\beta}}{\partial q}, \frac{\partial \hat{\beta}}{\partial \hat{p}_{human}} \right]^T = \left[\frac{1}{1 - \hat{p}_{human}}, \frac{q - 1}{(1 - \hat{p}_{human})^2} \right]^T \quad (7)$$

We can use this to compute the variance of $\hat{\beta}$:

$$\begin{aligned} Var(\hat{\beta}) &= Var(\hat{\beta}(\mu) + \nabla \hat{\beta}(\mu)^T ((Q, P) - \mu)) \\ &= Var(\hat{\beta}(\mu) + \nabla \hat{\beta}(\mu)^T (Q, P) - \nabla \hat{\beta}(\mu)^T \mu) \\ &= Var(\hat{\beta}(\mu)^T (Q, P)) \\ &= \nabla \hat{\beta}(\mu)^T \Sigma \nabla \hat{\beta}(\mu) \end{aligned} \quad (8)$$

We assume independence of the two random variables Q and P , so that the covariance matrix Σ only has the diagonal entries of the estimate variances SE_q^2

⁶As per common convention, the estimated regression coefficients are normally denoted $\hat{\beta}$, which was changed to $\hat{b}$ to avoid confusion with the LLM-use estimate $\hat{\beta}$

and $SE_{\hat{p}_{human}}^2$ derived above. We excluded all estimates with a standard error of $SE_{\hat{\beta}} \geq 0.025$.

Sanity check simulation

To better gauge how well the procedure of finding the optimal excess word list




Figure 5: Illustration of the sampling procedure for the creation of an artificial dataset to test the LLM use estimate $\hat{\beta}$ for varying levels of β

to get the best possible lower bound works, we performed the same analysis on simulated data with known p_{LLM} and p_{human} and varying β . We simulated data for 100,000 texts, approximating the number of papers per year in our sample. In our simulation, there were 500 words, all of which were treated as excess words. We created two matrices, M_{pre} and M_{post} , each with dimensions (*number of texts*, *number of words*), corresponding to data before and after the release of ChatGPT, respectively. For M_{pre} , each word w was randomly assigned a probability $p_{human}(w)$. The probabilities were sampled from a gamma distribution with shape parameter 2 and scale parameter 0.02. Then, 100,000 binomial samples were drawn with $p_{human}(w)$ as probability of success for each word. For M_{post} , a fraction of β texts were selected as being LLM-generated. As one of the assumptions posits that $p_{human} \leq p_{LLM}$, the probabilities for words to occur in LLM-assisted text were sampled as $p_{LLM}(w) = p_{human}(w)(1 + \delta)$, with $\delta \sim U(0.5, 5)$. In cases where $p_{human}(w)(1 + \delta) > 1$, δ was newly sampled until the resulting $p_{LLM}(w)$ was in the range $[0, 1]$. In the sampling process $100,000 \cdot \beta$ binomial samples were drawn with probability $p_{LLM}(w)$, and $100,000 \cdot (1 - \beta)$ were drawn with probability $p_{human}(w)$.

For the following analyses, the word frequencies in M_{pre} was treated as analogous to $\hat{p}_{human}$, and those in M_{post} as the observed frequencies q . The words in M_{pre} were ordered by frequency, and excess lists were formed of words with frequencies lower than those of several different cutoffs. For each excess list corresponding to one of the cutoffs, the LLM-use estimates Δ and $\hat{\beta}$ were computed. The only difference w.r.t. the procedure on the real data was during variance computation: M_{pre} , as opposed to $\hat{p}_{human}$, is not the result of a regression, but, similar to q the result of a binomial sampling procedure. Thus, the variance for frequencies based on M_{pre} was computed as the variance of a binomial estimator, $Var(M_{pre}) = \frac{p(1-p)}{n}$ with p the observed frequency and $n = (1 - \beta) \cdot n_{texts}$. As with the real data, we exclude cutoffs yielding frequencies

≥ 0.999 for M_{pre} , as well as estimates with error bars ≥ 0.025 .

3.5 Section length

One of the goals of this project was to compare the estimated LLM-use of different sections of the paper with each other. This might seem like an unfair comparison: When using word occurrence as LLM-marker, longer texts are at an advantage - there is simply more space for words to appear in. [Figure here shows the different section lengths] For example, if an LLM produced some word once every 500 words on average, it would appear in the majority of 1000-word sections, but only in about half of 250-word sections, despite all of them being written by an LLM. Since the word counts are binary, there is no correction for section length in the estimate. This is a further argument for using section specific word lists instead of single words. The longer the word list, and the higher their aggregated frequency, the less likely it is for an LLM-generated text, however short, to not contain at least one of them. However, when authors use LLMs, it is also likely that they focus their usage on specific paragraphs that they, e.g., wish to rephrase. A longer section contains more paragraphs that might warrant re-working. One way to address this issue is to slightly shift the question from *How high is LLM-use in this section?* to *How high is LLM-use in a random paragraph of this section?* This was implemented by cropping each section sample to a length of 255 words, which corresponds to the median length of the abstracts. The cropped part was placed randomly within the section by sampling a starting position ranging between the first and 255th-last words in the text. Note that the LLM-use estimate based on the cropped sections can't be used to determine the full fraction of papers that were written with some LLM participation. Since we are interested in identifying not only those papers that are wholly written by an LLM, using only one paragraph per section will underestimate the number of papers written with the help of LLMs, because this paragraph might not include the parts of the paper where LLMs were used. However, by comparing the results from the full and cropped sections, we can check whether LLMs are used universally across the entire paper or whether their use is concentrated on some paragraphs.

3.6 Subgroup analysis

Another point of interest is how LLM-use differs across subsets of the data. It is, for instance, plausible to assume that majority English-speaking countries will show different usage patterns than countries where English is learned as a second language. Likewise, there might be differences between academic disciplines, with some being quicker to adopt a new technology. Previous studies typically reported higher usage estimates for research areas related to computer science. To compare these categories, we repeated the experiments on different subsets of papers.

Countries

The most salient criterion to compare countries with each other when it comes to LLM-use is their majority native language. Since English is the de facto lingua franca of academia, one of the key capabilities of LLMs is their ability to translate or proofread text from authors uncertain of their English proficiency. As this use-case is irrelevant for native English speakers, it stands to reason that LLM-usage would be smaller for English-speaking countries.

To categorize countries into English and non-English speaking, we used the *Three circles of English model*[25], which sorts the United States, the United Kingdom and Ireland, Australia and New Zealand, South Africa, and the majority of Canada into the “inner circle”, where English is the native language of the majority of inhabitants and primarily used in their day-to-day lives. All other countries were classified as non-English speaking. Additionally to comparing English to non-English speaking countries, we also analyzed the LLM-use of individual countries. We selected the 18 countries with highest paper counts in the years 2017 to 2025. The full list of countries and their respective paper counts can be found in the Appendix.

We originally planned to use the excess word lists that are optimal for each section with regards to the full dataset. However, this lead to observations of $q = 1.0$ for several years and countries (pre- and post-release of ChatGPT). Due to time constraints, it was unfeasible to search for the optimal list for every country and section. As a heuristic to approximate the search process, the list corresponding to the next-lower cutoff was selected. For example, for the abstract, the optimal cutoff was 0.1, so we selected the list corresponding to a cutoff value of 0.07, which was the value directly preceding 0.1. By lowering the cutoff, we reduce the number of words in the excess word list and can thus regulate the observed frequencies. This “suboptimal” cutoff also lead to problematic observations, because in many cases, q was within the regression error of $\hat{p}_{human}$, which is not ideal for any metric depending on the difference between the two. To adjust, we moved the cutoff value one position further down, e.g. by going from 0.1 to 0.05 instead of 0.7 for the abstract.

Research disciplines

Papers were sorted into different research areas based on the titles of journals in which they were published in a procedure based on [18]. The research areas are subfields of the biomedical domain, for example radiology or nursing, that were specified in [18]. The full list of categories can be found in the Appendix. Journals were sorted into a category via keyword search with the category titles, and were assigned no label in case of no match. This is a relatively crude but simple approach that ignores many potential matches, e.g., the (hypothetical) *Journal of Neurology* would be assigned the label “neurology”, while the *Neurological Journal* would not. The number of matches could be improved by using fuzzy matching or manually annotating. In assigning labels, the keyword appearing latest in the title was given precedence. For example, the *Journal of Cancer Biochemistry* would be sorted into the “biochemistry” category, while *The Biochemistry of Cancer* would be assigned the label “cancer”. This procedure resulted in 23% of papers being labeled. Only labels with more than 5000

papers were selected for further analysis. Using the optimal excess word lists to compute frequencies lead to similar problems as detailed above, which is why we also used the cutoff twice removed from the optimal one.

4 Results

4.1 Simulation




Figure 6

4.2 Individual word estimates?

4.3 Word list estimates




Figure 7




Figure 8




Figure 9




Figure 10




Figure 11: Shared caption describing both images

5 Discussion

- advantage of not relying on ground-truth text: not vulnerable to different results stemming from different prompting strategies (Dou enhancing robustness 2024) (disadvantage is only lower bound, but as was shown, we get reasonably close to the actual value)
- going beyond word frequency: other linguistic markers as LLM indicators (perplexity, sentiment, part-of-speech, vocab size and density) e.g. [21], Bao Tong 2025, or intrinsic text embedding dimensionality [55], or repeated higher-order n-grams [14] (although this is for GPT2 and might be no longer true)
- non-english speaking word distribution starts to assimilate to english-speaking - ChatGPT as linguistic Equalizer (Lin 2025, Prakash 2025)
- impact / normative evaluation: overcoming language barriers, science papers rely on the data that is being reported, writing style is not immediately important. However, writing is a technology to sharpen our thinking. It takes effort to write something down. In deciding how exactly to phrase something, we are forced to think very precisely about our work, a level of scrutiny which makes it more likely that leaps of logic will be addressed that could not have been caught by the LLM

References

- [1] Arslan Akram. “Quantitative analysis of AI-generated texts in academic research: A study of AI presence in Arxiv submissions using AI detection tool”. In: *arxiv.org A Akram arXiv preprint arXiv:2403.13812, 2024* • *arxiv.org* (2024). This was not proofread. URL: <https://arxiv.org/abs/2403.13812>.
- [2] Aisha Alansari and Hamzah Luqman. “Large language models hallucination: A comprehensive survey”. In: *Computer Science Review* 61 (Aug. 2026), p. 100970. ISSN: 1574-0137. DOI: 10.1016/j.cosrev.2026.100970. URL: <http://arxiv.org/abs/2510.06265>.
- [3] Simone Astarita et al. “Delving into the Utilisation of ChatGPT in Scientific Publications in Astronomy”. In: (Sept. 2024). DOI: 10.5281/zenodo.13885576. URL: <http://arxiv.org/abs/2406.17324> %20http://dx.doi.org/10.5281/zenodo.13885576.
- [4] Yuntao Bai et al. “Training a Helpful and Harmless Assistant with Reinforcement Learning from Human Feedback”. In: (2022).
- [5] Tong Bao et al. “Examining linguistic shifts in academic writing before and after the launch of ChatGPT: a study on preprint papers”. In: *Scientometrics* 130 (7 July 2025), pp. 3597–3627. ISSN: 15882861. DOI: 10.1007/S11192-025-05341-Y/TABLES/9. URL: <https://link.springer.com/article/10.1007/s11192-025-05341-y>.
- [6] Violeta Berdejo-Espinola and Tatsuya Amano. “AI tools can improve equity in science”. In: *Science* 379 (6636 Mar. 2023), p. 991. ISSN: 10959203. DOI: 10.1126/SCIENCE.ADG9714. URL: <https://www.science.org/doi/10.1126/science.adg9714>.
- [7] Elouise Botes et al. “Initial indications of generative AI writing in linguistics research publications”. In: (2025). DOI: 10.31234/osf.io/4yvbvp_v1. URL: https://www.researchgate.net/profile/Elouise-Botes/publication/394388056%5C_Initial%5C_indications%5C_of%5C_generative%5C_AI%5C_writing%5C_in%5C_linguistics%5C_research%5C_publications/links/68a5e6fe6327cf7b63d83be1/Initial-indications-of-generative-AI-writing-in-linguistics-research-publications.pdf.
- [8] Jeffrey Brainard. “Journals take up arms against AI-written text”. In: *Science* 379 (6634 Feb. 2023), pp. 740–741. ISSN: 10959203. DOI: 10.1126/SCIENCE.ADH2762.
- [9] Hu-Zi Cheng et al. “Have AI-generated texts from LLM infiltrated the realm of scientific writing? A large-scale analysis of preprint platforms”. In: *biorxiv.org HZ Cheng, B Sheng, A Lee, V Chaudhary, AG Atanasov, N Liu, Y Qiu, TY Wong, YC Tham bioRxiv, 2024* • *biorxiv.org* (2024). DOI: 10.1101/2024.03.25.586710. URL: <https://www.biorxiv.org/content/10.1101/2024.03.25.586710.abstract>.

- [10] Heather Desaire et al. “Distinguishing academic science writing from humans or ChatGPT with over 99% accuracy using off-the-shelf machine learning tools”. In: *Cell Reports Physical Science* 4 (6 June 2023), p. 101426. ISSN: 2666-3864. DOI: 10.1016/J.XCRP.2023.101426.
- [11] Jad Doughman et al. *Exploring the Limitations of Detecting Machine-Generated Text*. 2025. URL: <https://aclanthology.org/2025.coling-main.288/>.
- [12] Elsevier. *Insights 2024 — Attitudes toward AI — Elsevier*. 2024. URL: <https://www.elsevier.com/insights/attitudes-toward-ai>.
- [13] Michael Eppler et al. “Awareness and Use of ChatGPT and Large Language Models: A Prospective Cross-sectional Global Survey in Urology”. In: *European Urology* 85 (2 Feb. 2024), pp. 146–153. ISSN: 0302-2838. DOI: 10.1016/J.EURURO.2023.10.014.
- [14] Matthias Gallé et al. “Unsupervised and Distributional Detection of Machine-Generated Text”. In: (Nov. 2021). URL: <http://arxiv.org/abs/2111.02878>.
- [15] Sebastian Gehrmann, Hendrik Strobelt, and Alexander M. Rush. “GLTR: Statistical Detection and Visualization of Generated Text”. In: *ACL 2019 - 57th Annual Meeting of the Association for Computational Linguistics, Proceedings of System Demonstrations* (2019), pp. 111–116. DOI: 10.18653/V1/P19-3019. URL: <https://aclanthology.org/P19-3019/>.
- [16] Mingmeng Geng and Roberto Trotta. “Human-LLM Coevolution: Evidence from Academic Writing”. In: (Feb. 2025), pp. 12689–12696. DOI: 10.18653/v1/2025.findings-acl.657. URL: <https://arxiv.org/abs/2502.09606v2>.
- [17] Mingmeng Geng and Roberto Trotta. “Is ChatGPT Transforming Academics’ Writing Style?” In: *arxiv.orgM Geng, R TrottaarXiv preprint arXiv:2404.08627, 2024•arxiv.org* (2024). URL: <https://arxiv.org/abs/2404.08627>.
- [18] Rita González-Márquez et al. “The landscape of biomedical research”. In: *Patterns* 5.6 (2024).
- [19] Andrew Gray. “ChatGPT” contamination”: estimating the prevalence of LLMs in the scholarly literature”. In: *arxiv.orgA GrayarXiv preprint arXiv:2403.16887, 2024•arxiv.org* 53 (6 2024), pp. 331–333. DOI: 10.1002/leap.1578. URL: <https://arxiv.org/abs/2403.16887>.
- [20] Andrew Gray. “Estimating the prevalence of LLM-assisted text in scholarly writing”. In: *arxiv.org* 53 (6 Dec. 2025), pp. 331–333. DOI: 10.1002/leap.1578. URL: <https://arxiv.org/abs/2512.01560>.
- [21] Biyang Guo et al. “How Close is ChatGPT to Human Experts? Comparison Corpus, Evaluation, and Detection”. In: (Jan. 2023). URL: <https://arxiv.org/abs/2301.07597v1>.

- [22] Abhimanyu Hans et al. “Spotting LLMs With Binoculars: Zero-Shot Detection of Machine-Generated Text”. In: *Proceedings of Machine Learning Research* 235 (Jan. 2024), pp. 17519–17537. ISSN: 26403498. URL: <https://arxiv.org/abs/2401.12070v3>.
- [23] Alex Hern. *TechScape: How cheap, outsourced labour in Africa is shaping AI English — Technology — The Guardian*. Apr. 2024. URL: <https://www.theguardian.com/technology/2024/apr/16/techscape-ai-gadgest-humane-ai-pin-chatgpt>.
- [24] Tom S. Juzek and Zina B. Ward. “Why Does ChatGPT “Delve” So Much? Exploring the Sources of Lexical Overrepresentation in Large Language Models”. In: *Proceedings - International Conference on Computational Linguistics, COLING* (2025), pp. 6397–6411. ISSN: 29512093.
- [25] Braj B Kachru. “Standards, codification and sociolinguistic realism: The English language in the outer circle”. In: *English in the world: Teaching and learning the language and literatures/Cambridge UP* (1985).
- [26] Graham Kendall and Jaime A Teixeira Da Silva. “Risks of abuse of large language models, like ChatGPT, in scientific publishing: Authorship, predatory publishing, and paper mills”. In: *Learned Publishing* 37 (2024). DOI: 10.1002/leap.1578. URL: <https://onlinelibrary.wiley.com/doi/10.1002/leap.1578>.
- [27] Dmitry Kobak et al. “Delving into LLM-assisted writing in biomedical publications through excess vocabulary”. In: *Science Advances* 11 (27 July 2025), p. 3813. ISSN: 23752548. DOI: 10.1126/SCIADV.ADT3813/SUPPL\ _FILE/SCIADV.ADT3813_SM.PDF. URL: <https://www.science.org/doi/10.1126/sciadv.adt3813>.
- [28] Kayvan Kousha. “How is ChatGPT acknowledged in academic publications?” In: *Scientometrics* 129 (12 Dec. 2024), pp. 7959–7969. ISSN: 15882861. DOI: 10.1007/S11192-024-05193-Y/FIGURES/4. URL: <https://link.springer.com/article/10.1007/s11192-024-05193-y>.
- [29] Kayvan Kousha and Mike Thelwall. “How much are LLMs changing the language of academic papers after ChatGPT? A multi-database and full text analysis”. In: *Scientometrics* 2026 (Sept. 2025), pp. 1–21. ISSN: 15882861. DOI: 10.1007/S11192-026-05601-5/FIGURES/14. URL: <http://arxiv.org/abs/2509.09596>.
- [30] Diana Kwon. “Is it OK for AI to write science papers? Nature survey shows researchers are split”. In: *Nature* 641 (8063 May 2025), pp. 574–578. ISSN: 14764687. DOI: 10.1038/D41586-025-01463-8.
- [31] Teddy Lazebnik and Ariel Rosenfeld. “Detecting LLM-assisted writing in scientific communication: Are we there yet?” In: *JDIS Journal of Data and Information Science* (2024). DOI: 10.2478/jdis-2024-0020. URL: <https://sciendo.com/journal/JDIS>.

- [32] Weixin Liang et al. “GPT detectors are biased against non-native English writers”. In: *Patterns* 4 (7 July 2023), p. 100779. ISSN: 26663899. DOI: 10.1016/j.patter.2023.100779. URL: [https://www.cell.com/action/showFullText?pii=S2666389923001307%20https://www.cell.com/action/showAbstract?pii=S2666389923001307%20https://www.cell.com/patterns/abstract/S2666-3899\(23\)00130-7](https://www.cell.com/action/showFullText?pii=S2666389923001307%20https://www.cell.com/action/showAbstract?pii=S2666389923001307%20https://www.cell.com/patterns/abstract/S2666-3899(23)00130-7).
- [33] Weixin Liang et al. “Mapping the increasing use of LLMs in scientific papers”. In: *arxiv.org* W Liang, Y Zhang, Z Wu, H Lepp, W Ji, X Zhao, H Cao, S Liu, S He, Z Huang, D YangarXiv preprint arXiv:2404.01268, 2024•arxiv.org (2024). URL: <https://arxiv.org/abs/2404.01268>.
- [34] Weixin Liang et al. “Monitoring ai-modified content at scale: A case study on the impact of chatgpt on ai conference peer reviews”. In: *arxiv.org* W Liang, Z Izzo, Y Zhang, H Lepp, H Cao, X Zhao, L Chen, H Ye, S Liu, Z HuangarXiv preprint arXiv:2403.07183, 2024•arxiv.org (2024). URL: <https://arxiv.org/abs/2403.07183>.
- [35] Weixin Liang et al. “Quantifying large language model usage in scientific papers”. In: *Nature Human Behaviour* 2025 9:12 9 (12 Aug. 2025), pp. 2599–2609. ISSN: 2397-3374. DOI: 10.1038/s41562-025-02273-8. URL: <https://www.nature.com/articles/s41562-025-02273-8>.
- [36] Zhehui Liao et al. “LLMs as Research Tools: A Large Scale Survey of Researchers’ Usage and Perceptions”. In: 1 (Oct. 2024). DOI: 10.1145/nnnnnnn.nnnnnnn. URL: <https://arxiv.org/abs/2411.05025v1>.
- [37] Michael Liebrecht et al. “Generating scholarly content with ChatGPT: ethical challenges for medical publishing”. In: *The Lancet Digital Health* 5 (3 Mar. 2023), e105–e106. ISSN: 25897500. DOI: 10.1016/S2589-7500(23)00019-5. URL: [https://www.thelancet.com/action/showFullText?pii=S2589750023000195%20https://www.thelancet.com/action/showAbstract?pii=S2589750023000195%20https://www.thelancet.com/journals/landig/article/PIIS2589-7500\(23\)00019-5/abstract](https://www.thelancet.com/action/showFullText?pii=S2589750023000195%20https://www.thelancet.com/action/showAbstract?pii=S2589750023000195%20https://www.thelancet.com/journals/landig/article/PIIS2589-7500(23)00019-5/abstract).
- [38] Zhicheng Lin. “Hidden Prompts in Manuscripts Exploit AI-Assisted Peer Review”. In: (July 2025). URL: <https://arxiv.org/abs/2507.06185v1>.
- [39] Grace W. Lindsay. “LLMs are not ready for editorial work”. In: *Nature Human Behaviour* 2023 7:11 7 (11 Nov. 2023), pp. 1814–1815. ISSN: 2397-3374. DOI: 10.1038/s41562-023-01730-6. URL: <https://www.nature.com/articles/s41562-023-01730-6>.
- [40] Jialin Liu and Yi Bu. “Towards the relationship between AIGC in manuscript writing and author profiles: evidence from preprints in LLMs”. In: (Apr. 2024). URL: <http://arxiv.org/abs/2404.15799>.
- [41] Miles MacClure. *Digital Frontier — Our new AI dialect*. Oct. 2024. URL: <https://digitalfrontier.com/articles/ai-dialects-delve-chatgpt>.

- [42] Matthew HC Mak and Lukasz Walasek. “Style, Sentiment, and Quality of Undergraduate Writing in the AI Era: A Cross-Sectional and Longitudinal Analysis of 4,820 Authentic Empirical Reports”. In: (Aug. 2025). DOI: 10.21203/RS.3.RS-7465165/V1. URL: <https://www.researchsquare.com/https://www.researchsquare.com/article/rs-7465165/v1>.
- [43] Kentaro Matsui. “Delving into PubMed records: Some terms in medical writing have drastically changed after the arrival of ChatGPT”. In: *medrxiv.orgK MatsuiMedRxiv, 2024*•*medrxiv.org* (2024). DOI: 10.1101/2024.05.14.24307373. URL: <https://www.medrxiv.org/content/10.1101/2024.05.14.24307373.abstract>.
- [44] National Library of Medicine. *PMC Open Access Subset [Internet]*. <https://pmc.ncbi.nlm.nih.gov/tools/openftlist/>. [cited 2026 07 04]. Bethesda (MD), 2003.
- [45] Tanisha Mishra et al. “Use of large language models as artificial intelligence tools in academic research and publishing among global clinical researchers”. In: *Scientific Reports* 14 (2024), p. 31672. DOI: 10.1038/s41598-024-81370-6. URL: <https://www.nature.com/articles/s41598-024-81370-6>.
- [46] Eric Mitchell et al. *DetectGPT: Zero-Shot Machine-Generated Text Detection using Probability Curvature*. July 2023. URL: <https://proceedings.mlr.press/v202/mitchell123a.html>.
- [47] Ehsan Mohammadi et al. “Is generative AI reshaping academic practices worldwide? A survey of adoption, benefits, and concerns”. In: *Information Processing & Management* 63 (1 Jan. 2026), p. 104350. ISSN: 0306-4573. DOI: 10.1016/J.IPM.2025.104350.
- [48] Richard Van Noorden and Jeffrey M. Perkel. “AI and science: what 1,600 researchers think”. In: *Nature* 621 (7980 Sept. 2023), pp. 672–675. ISSN: 14764687. DOI: 10.1038/D41586-023-02980-0.
- [49] Linda Nordling. “How ChatGPT is transforming the postdoc experience”. In: *Nature* 622 (7983 Oct. 2023), pp. 655–657. ISSN: 14764687. DOI: 10.1038/D41586-023-03235-8. URL: <https://www.nature.com/articles/d41586-023-03235-8>.
- [50] Billy Perrigo. *OpenAI Used Kenyan Workers on Less Than \$2 Per Hour: Exclusive*. Jan. 2023. URL: <https://time.com/6247678/openai-chatgpt-kenya-workers/>.
- [51] Pablo Picazo-Sanchez and Lara Ortiz-Martin. “Analysing the impact of ChatGPT in research”. In: *Applied Intelligence* 54 (5 Mar. 2024), pp. 4172–4188. ISSN: 15737497. DOI: 10.1007/S10489-024-05298-0/TABLES/8. URL: <https://link.springer.com/article/10.1007/s10489-024-05298-0>.

- [52] Arjun Prakash et al. “Writing without borders: AI and cross-cultural convergence in academic writing quality”. In: *Humanities and Social Sciences Communications* 2025 12:1 12 (1 July 2025), pp. 1058–. ISSN: 2662-9992. DOI: 10.1057/s41599-025-05484-6. URL: <https://www.nature.com/articles/s41599-025-05484-6>.
- [53] Vinu Sankar Sadasivan et al. “Can AI-generated text be reliably detected?” In: *arxiv.org* (Jan. 2025). URL: <https://arxiv.org/abs/2303.11156>.
- [54] H. Holden Thorp. “ChatGPT is fun, but not an author”. In: *Science* 379 (6630 Jan. 2023), p. 313. ISSN: 10959203. DOI: 10.1126/SCIENCE.ADG7879/ASSET/E576EF79-5CF7-44F0-A6F6-8AA953COD70F/ASSETS/IMAGES/LARGE/SCIENCE.ADG7879-F1.JPG. URL: <https://www.science.org/doi/10.1126/science.adg7879>.
- [55] Eduard Tulchinskii et al. “Intrinsic Dimension Estimation for Robust Detection of AI-Generated Texts”. In: *37th Conference on Neural Information Processing Systems (NeurIPS 2023)* (2023).
- [56] Sergio E. Uribe and Ilze Maldupa. “Estimating the use of ChatGPT in dental research publications”. In: *Journal of Dentistry* 149 (Oct. 2024), p. 105275. ISSN: 0300-5712. DOI: 10.1016/J.JDENT.2024.105275.
- [57] William H. Walters and Esther Isabelle Wilder. “Fabrication and errors in the bibliographic citations generated by ChatGPT”. In: *Scientific Reports* 2023 13:1 13 (1 Sept. 2023), pp. 14045–. ISSN: 2045-2322. DOI: 10.1038/s41598-023-41032-5. URL: <https://www.nature.com/articles/s41598-023-41032-5>.
- [58] Debora Weber-Wulff et al. “Testing of detection tools for AI-generated text”. In: *International Journal for Educational Integrity* 19 (1 Dec. 2023). ISSN: 18332595. DOI: 10.1007/S40979-023-00146-Z.
- [59] Wiley. *ExplanAltions: An AI study by Wiley*. Feb. 2025. URL: <https://www.wiley.com/content/dam/wiley-com/en/pdfs/about/wiley-explanAltions-ai-study-february-2025vers3.pdf>.
- [60] Ziwei Xu, Sanjay Jain, and Mohan Kankanhalli. “Hallucination is Inevitable: An Innate Limitation of Large Language Models”. In: (Feb. 2025). URL: <http://arxiv.org/abs/2401.11817>.
- [61] Haoyi Zheng and Huichun Zhan. “ChatGPT in Scientific Writing: A Cautionary Tale”. In: *American Journal of Medicine* 136 (8 Aug. 2023), 725–726.e6. ISSN: 15557162. DOI: 10.1016/j.amjmed.2023.02.011. URL: [https://www.amjmed.com/action/showFullText?pii=S0002934323001596%20https://www.amjmed.com/action/showAbstract?pii=S0002934323001596%20https://www.amjmed.com/article/S0002-9343\(23\)00159-6/abstract](https://www.amjmed.com/action/showFullText?pii=S0002934323001596%20https://www.amjmed.com/action/showAbstract?pii=S0002934323001596%20https://www.amjmed.com/article/S0002-9343(23)00159-6/abstract).

A Survey results

In all surveys, we focused on use cases that could have an influence on the formulation of final paper, such as editing or summarizing of other papers, instead of use cases in other academic undertakings, such as writing code, or very early in the writing process, like brainstorming or writing an outline. Grammar checks were also excluded as a relevant use case, because while this is a very popular task for LLMs, this should have a minor influence on phrasing. If use cases were not part of the survey, we took the general proportion of researchers reporting LLM use in at least some part of their research. an important factor in evaluating the results is the organization doing the survey. Researchers might be hesitant to admit LLM use in a survey done by an important academic publisher despite anonymity, especially if this publisher is known to be anti-LLM, such as Nature

survey period	sample size	sample specification	question	category	usage	source
04/2023–05/2023	456	urologists	Do you use ChatGPT/LLMs in your academic practice? summarizing text	general writing	48% 27%	[13]
06/2023–07/2023	3,838	postdocs	Do you use AI chatbots, such as ChatGPT, in your work? What do you use AI chatbots for? Refining text	general editing	31% 20%	[49]
?–09/2023	1,659	researchers who had published in Q4 2022	What do you use generative AI tools (e.g., ChatGPT, LLMs) for? To help write research manuscripts	writing	21%	[48]
11/2023–04/2024	816	researchers (Semantic Scholar)	Fix grammar or rephrase Rewrite for another style Shorten or summarize Draft paragraphs from ideas Direct writing (overall)	editing writing writing writing writing	50% 24% 23% 21% 43%	[36]
12/2023–02/2024	2,284	clinicians & researchers	Have you used an AI product or AI feature for a specific work-related purpose?	general	37%	[12]
01/2024–03/2024	2,021	researchers (Scopus)	Why use ChatGPT/AI tools: Synthesizing drafts to academic text Proofreading drafts Translating text	writing editing writing	16% 17% 18%	[47]
03/2024–04/2024	1,043	researchers	Help with translation proofreading and editing scholarly papers	writing editing	40% 38%	[59]
03/2024–05/2024	5,229	researchers	Have you used AI to edit your paper? ... to translate your paper? ... to summarize other papers for use in your own? ... to write the first draft of your paper?	editing writing writing writing	28% 8% 8% 8%	[30]
04/2024–06/2024	226	medical researchers	Stage of publication LLM used: Writing Revision and editing	writing editing	8% 8%	[45]

Table 1: Summary of survey results on AI/LLM usage in academic and clinical research

B 2024 excess words

The following is a list of all excess style words identified by [27]. They are sorted from least to most frequent, as measured for abstracts in PMC.

revolves, upholding, dependability, attains, endeavours, scrutinizing, garnering, consolidates, encapsulates, merges, accentuates, grappling, acknowledges, furnish, delving, solidify, uphold, urging, surmount, juxtaposed, customizing, propelling, detrimentally, interconnectedness, equipping, realms, categorizes, commendable, embracing, crafting, substantiates, bolstered, pioneers, burgeoning, unlocking, capitalizing, heighten, gauged, emulating, refines, uncharted, unparalleled, crafted, adept, scrutinize, swiftly, emulation, excels, surged, harnesses, revolutionizing, spurred, streamlines, expediting, dependable, bolstering, empowers, groundbreaking, illuminating, hinting, discerning, afflicted, inquiries, scrutinized, detailing, discerned, surpass, overlook, streamlining, emphasising, delineates, seamless, overlooking, seamlessly, hinges, emphasises, acknowledging, revolutionize, fosters, pinpointing, nuances, swift, advocating, meticulous, strategically, signifying, delved, illuminates, surpasses, amidst, renowned, outlines, complicating, spanned, showcased, intricacies, advocates, adhering, bolster, categorizing, poised, stemming, delve, expedite, formidable, affirming, complicates, pioneering, endeavors, culminating, elevates, delves, adhered, intricately, comprehending, showcases, substantiated, streamline, transformative, unveiling, meticulously, streamlined, harness, realm, surpassed, enduring, pinpointed, distinctions, impressive, imbalances, inadequately, exceptionally, necessitate, discernible, elevating, elevate, exerting, invaluable, struggle, enhancements, showcasing, aligns, encompass, diminishing, unraveling, harnessing, discern, diminishes, akin, garnered, refining, impeding, escalating, exacerbating, augmenting, orchestrating, impede, warranting, pinpoint, foundational, groundwork, seeks, underscored, encompasses, pressing, combating, versatility, unveils, surpassing, aligning, posed, paving, fostering, featuring, noteworthy, elucidates, necessitates, stands, uncovering, aiding, posing, leverages, hinder, exceptional, encompassed, nuanced, compelling, refine, incorporates, underexplored, thorough, employs, align, recognizing, evaluates, methodologies, correlating, addresses, unveiled, utilizes, unveil, necessitating, multifaceted, imperative, exceeding, emerges, evolving, displaying, integrates, yielding, avenue, disrupts, introduces, examines, advancement, shedding, emphasizes, elucidating, preserving, persist, distinctive, inherent, necessity, serving, advancements, optimizing, advancing, explores, spanning, underscoring, shaping, avenues, hold, impacting, emphasize, ensuring, encompassing, pose, emphasizing, heightened, remarkable, unexplored, leveraging, intricate, investigates, broader, alongside, serves, holds, poses, capabilities, facilitated, facilitates, innovative, employing, initially, underscores, exploration, exhibiting, incorporating, attributed, interplay, pronounced, involves, addressing, comprising, revealing, notable, ultimately, integrating, categorized, underscore, facilitating, conversely, achieving, focusing, consequently, promise, exhibits, predominantly, influencing, enabling, presents, integration, offering, typically, demonstrating, varying, assessments, enhances,

maintaining, pivotal, utilizing, contributing, precise, limitations, elucidate, emerged, offers, assessing, highlighting, offer, reveals, prevalent, involving, techniques, subsequently, primarily, enhancing, demonstrates, utilized, substantial, exhibit, challenge, thereby, valuable, challenges, effectively, highlights, address, highlight, employed, providing, enhance, specifically, indicating, diverse, notably, comprehensive, resulting, potentially, elevated, linked, promising, leading, aims, enhanced, distinct, particularly, despite, need, strategies, insights, crucial, exhibited, like, conditions, complex, additionally, individuals, various, assess, understanding, approach, remains, primary, assessed, impact, outcomes, demonstrated, conducted, across, research, within, revealed, observed, through, role, identified, while, including, into, potential, findings, both, their, using, these, this

C Additional simulation results




Figure 12




Figure 13

D Word list optimization




Figure 14




Figure 15